

Bloom Filter and Variation


Shicheng Fang

Bloom Filter: What is it?

- Test an element exists in set
- Support two operations:
 - **add(key)**
 - **isExist(key)**
- $O(1)$ Time and Space complexity
- Possible erroneous result for **isExist(key)**
 - **isExist(key)** returns true though key was not in the set
 - We call it **False Positive**


Bloom Filter: How does it work?

- Initially, an empty bit array filled with 0 of size M
- M is usually a large integer




Bloom Filter: How does it work? (continue)

- Add(key)
 - K hash functions calculate K values as array index
 - The K array indexes will be filled with 1
- Example:
 - $K=3$
 - $\text{hash1}(\text{key})=0$, $\text{hash2}(\text{key})=3$, $\text{hash3}=6$



Bloom Filter: How does it work? (continue)

- `isExist(key)`
 - K hash functions calculate K values as array index
 - Return if all K locations are 1
- Example:
 - $K=3$
 - $\text{hash1}(\text{key})=0$, $\text{hash2}(\text{key})=3$, $\text{hash3}=6$
 - $\text{Array}[6] \neq 1$, return false



Bloom Filter: False Positive

- `isExist(key)` returns true even if key was not added before
- Consider Bloom Filter filled with 1
 - `isExist(key)` ALWAYS returns true, not useful at all

Bloom Filter: Tradeoff

- Tradeoff between:
 - False Positive rate (fp)
 - Bloom filter size (M)
 - Number of items in the filter (N)
 - Number of hash functions (K)
- In general,
 - $K = \ln(2) * M/N$ round to integer
 - Higher M => Lower fp
 - Higher N => High fp

Bloom Filter: Performance Tuning

- Calculate index using bitwise AND instead of %
 - Bitwise AND is faster than %
 - $A \% B == A \& (\text{LOG2}(B) - 1)$
 - B is power of 2
- Calculate hash for once instead of k
 - $\text{Base, Inc} = \text{hash_function}(\text{key})$
 - Set $\text{Base} + 1 * \text{Inc}$, $\text{Base} + 2 * \text{inc}$, ..., $\text{Base} + k * \text{inc}$

What if M is not power of 2?

- Enforce M to be power of 2 (Inflexible)
- Use % (Low performance)
- Turn % into a sequence of faster operations
 - Usually, sequence is Multiply + Shift + Add/Sub

Multiply, Shift, Add

- Common technique in compiler optimization of %
- However, M is unknown at compilation time
 - Compiler cannot optimize %
- Solution:
 - Make M available at compilation time
 - Need recompilation when M changes
 - Very inflexible
 - Perform optimized % at runtime

Multiply, Shift, Add

- $A \% B = A - (A / B) * B$
- A / B takes most clock cycles
 - How to optimize division?

Faster division

E.g. Divide by 3

```
computeDivideBy3(long x) {  
    long y = x;  
    y = y * 1431655766;  
    y = y >> 32;  
    x = x >> 31;  
    return y - x;  
}
```

- Technique called **division by reciprocal approximation**
- 1431655766 is magic number.
- **EVERY** divisor has its own magic number.

What is 1431655766?

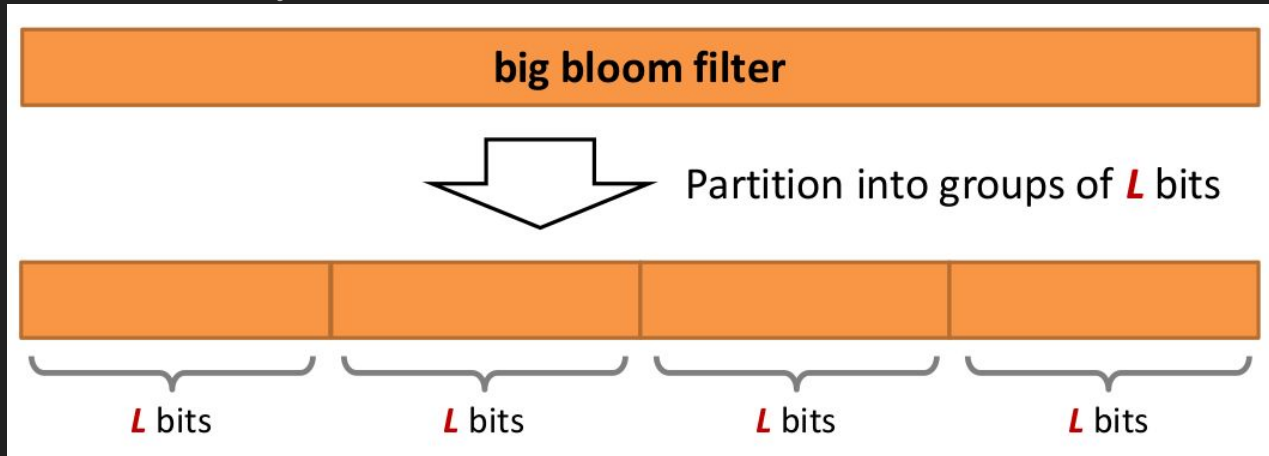
- Convert the divisor 3 to its reciprocal, which is $1/3$.
- Multiply the reciprocal by a large power of 2, e.g. 2^{30}
 - $1/3 * 2^{30} = 1073741824 / 3$
- Round the resulting value to the nearest integer, which gives us 1431655766.
- This value is precalculated for once and stored for later use.

MOD or Multiply, shift, add

- Which is faster? (On Pentium CPU)
 - DIV takes 40 clock cycles.
 - MUL takes 10 clock cycles.
 - Add/Sub takes 1 clock cycle
 - Multiply, shift, add is significantly faster than MOD
- On other platform, MOD may be faster

Bloom Filter: Blocked Filter Variation

- Observation:
 - Set K bits can cause at most K cache miss
- Block Filter:
 - Divide the array into multiple cache lines each with size L bits, $L=512$
 - Set k bits in one cache line
 - Only 1 cache miss



Source:
Design Tradeoffs of
Bloom Filters
Jianguo Wang (UCSD)

Blocked Bloom Filter Disadvantage

- Higher FP rate (Bits are squeezed into a cache line)
- K bits should be set separately (K writes are needed)

Bloom Filter: Register Bloom Filter Variation

- Observation:
 - Set K bits need K write (K MOV instructions)
- Register Filter:
 - Divide the array into blocks with size of a register (64 bits) rather than blocks (512 bits)
 - K bits are set simultaneously by constructing the mask and then doing OR :
 - E.g. to set index 0, 3, 5, and 6, only OR 101011
 - MOV is executed for once

Register Bloom Filter Disadvantage


- Even higher FP (K bits are squeezed in 64 bits)

Bloom Filter: Sectorized Filter Variation

- Blocked Bloom Filter: too slow
- Register Bloom Filter: high fp
- Can we make a tradeoff?
 - Sectorized Filter

Bloom Filter: Sectorized Filter Variation

- An extension of Blocked Bloom Filter
- One block is divided into multiple into sectors (8 sectors as shown)
- To set K bits, K bits are set in C sectors




Bloom Filter: Sectorized Filter Variation (cont.)

Example:

8 sector, each with size 64 bits

$C=2$, $K=8$




Bloom Filter: Sectorized Filter Variation (cont.)

Example: (cont.)

$K=8$, $C=2$, meaning 2 sectors are used, e.g. $r0$ and $r1$.

4 bits set in each sector, 4 bits can be set simultaneously like a register block bloom filter





Performance Comparison

STD is the standard bloom filter used in Cassandra and RockDB

$n = 147639508$

$m = 1181116006$ bits (too large, stored in L3/dram)

$k = 2$



Small size favors register blocking

$n = 132416$

$m = 1059061$ bits (L1)

K:8

